

Full baseline sweep: best-representation macro-F1 by task, model, and CV split (weak-label agreement, Section 4.3)

