

---

# **RansKnow: An Open Dataset of Ransomware-Related YouTube Transcripts with MITRE ATT&CK Annotations**

---

**Henry Kabuye**

School of Computer Science  
Northumbria University, Newcastle upon Tyne, NE1 8ST, UK  
h.kabuye@northumbria.ac.uk

**Biju Issac**

School of Computer Science  
Northumbria University, Newcastle upon Tyne, NE1 8ST, UK  
bissac@ieee.org

## **Abstract**

We present RansKnow, an open corpus of 1,128 ransomware-related YouTube video records (1,034 with extracted transcripts) drawn from 92 cybersecurity channels, with MITRE ATT&CK-aligned structured features – ransomware family, tactic, tool, and platform mentions – produced by a deterministic regex-based extraction pipeline. This paper’s contribution is not the corpus alone: a systematic label-quality audit found that 9 of 41 ransomware-family aliases collide with ordinary English or business words (e.g. play, hive, cactus) badly enough to inflate the single most-detected family by roughly  $3.5\times$  before a context-based disambiguation fix, and a counterfactual-explanation analysis later in the paper shows a trained classifier partially inherited that same word-level fragility even after the fix. We state plainly here, not only in a limitations section, that this project does not include an independently human-annotated evaluation set: every baseline, calibration, and conformal-prediction result we report is agreement with this audited-but-imperfect extraction pipeline, not verified real-world accuracy, and we frame every such number accordingly throughout. Under that scope, weak-label baselines reach macro-F1 0.90 on a binary relevance screen and 0.81 on dominant-tactic classification, but remain weak (macro-F1  $\leq 0.20$ ) on the higher-cardinality family and tool taxonomies; calibration against the weak labels is moderate (Expected Calibration Error 0.02–0.11) with a Random Forest tree-vote-entropy signal that separates label agreement from disagreement well (AUROC up to 0.98); and split conformal prediction achieves close to its 90% target coverage with no severe conditional-coverage collapse on two out-of-calibration-distribution slices tested. Out-of-distribution detection, which does not depend on label correctness, tells a more cautionary story: standard softmax- confidence-based methods perform near chance across four realistic distribution-shift setups, including a diagnosed domain-specific failure mode where the classifier’s most confident predictions coincide with exactly the signature that off-topic input also produces; only embedding-based Mahalanobis distance clears chance, and only on one of the four setups. We release the corpus, the full extraction and audit pipeline, and all experiment code openly, and argue that the audit methodology – not just its findings – is reusable by any weak-label security-NLP dataset built the same way.

# 1 Introduction

Ransomware knowledge is scattered across the same kind of source material that generated it in the first place: incident-response writeups, conference talks, and vendor threat-intelligence reporting, almost none of it structured for downstream NLP or ML use. Structured resources do exist – the MITRE ATT&CK framework itself [Strom et al., 2018], and datasets built against it from written CTI text (Section 2) – but no open, transcript-level corpus covers the specific and large body of ransomware knowledge that security professionals communicate *by speaking*: conference talks, vendor webinars, and independent-researcher walkthroughs published as video. YouTube specifically is a useful aggregation point for this material precisely because of its source diversity: DFIR vendors publishing incident retrospectives, security conferences publishing full talk recordings, and independent researchers publishing technical walkthroughs all use the same platform, under different incentives and house styles, in a way no single vendor blog or CTI feed replicates. That diversity is also why transcribing it is nontrivial – 92.7% of this corpus required automatic speech recognition rather than official captions (Section 3.3), which turns out to matter for downstream label quality, not just as a data-engineering footnote (Section 4.1).

This paper makes five contributions. First, RansKnow itself: 1,128 video records, 1,034 with extracted transcripts, from 92 cybersecurity YouTube channels, with MITRE ATT&CK-aligned structured features (Section 3). Second, a label-quality audit methodology and its concrete findings (Section 4): a systematic, dictionary-driven search for generic-word alias collisions in the regex-based extraction pipeline, which found that 9 of 41 family aliases collide with ordinary English or business words badly enough to inflate the single most-detected family by roughly  $3.5\times$  before a disambiguation fix (Table 1). Third, an experiment framework and weak-label baseline results across five prediction tasks, three feature representations, and five models (Section 5), reported throughout – here in the introduction, not only in Section 10 where a reader could miss it – as agreement with the audited extraction pipeline, *not* as independently verified accuracy: this project does not include a human-annotated evaluation set held out from the labeling pipeline (Section 4.3), and every quantitative claim from Section 5 through Section 7 should be read with that scope in mind. Fourth, out-of-distribution robustness results (Section 8) that are explicitly *not* subject to that caveat, since they measure whether a model’s own confidence or embedding behavior differs on distributionally shifted input, which is answerable without knowing whether any label is correct. Fifth, counterfactual and perturbation-robustness interpretability results (Section 9) that connect back to the audit: a trained classifier is shown to have partially inherited the same word-level fragility Finding A identifies in the labeling pipeline itself, evidence that the audit’s concerns are not confined to the raw labels but propagate into what gets learned from them.

## 2 Related Work

**Ransomware and cyber-threat-intelligence datasets.** The MITRE ATT&CK framework [Strom et al., 2018] that RansKnow’s tactic taxonomy is aligned to was designed for exactly this kind of cross-source mapping, and prior datasets have built structured resources against it from written cyber-threat-intelligence (CTI) text: MalwareTextDB [Lim et al., 2017] provides span-level annotations (39 APT reports, 6,819 sentences) for malware characteristics, and the more recent AnnoCTR [Lange et al., 2024] links entities, tactics, and techniques across 400 commercial CTI reports, with human annotation throughout. Both are built from written prose – vendor blog posts and reports, produced already in structured-adjacent form for a threat-intelligence-consuming audience. RansKnow’s source material is instead spoken narration, transcribed at 92.7% via ASR (Section 3.3), from a broader and more heterogeneous set of speakers (vendor analysts, conference presenters, independent researchers) than a CTI report’s house style, and its structured features are produced entirely by an audited regex/keyword pipeline (Section 4) rather than human annotation – a different, explicitly weaker, point on the label-quality spectrum those two datasets occupy, traded for scale and openness (no comparable open, transcript-level, MITRE-aligned video corpus exists at the time of writing).

**Weak and distant supervision.** Using an automated heuristic in place of human labels is a well-studied tradeoff, not a shortcut unique to this project: Mintz et al. [2009] formalized distant supervision for relation extraction by aligning text against a knowledge base rather than hand-labeling it, accepting a known noise ceiling in exchange for scale, and Ratner et al. [2017]’s Snorkel generalizes the idea to multiple, possibly conflicting, weak labeling functions combined probabilistically. The

Knowledge Agent (Section 4) is a single deterministic labeling function rather than an ensemble of weak sources, which makes it simpler to audit (the alias-collision analysis in Section 4.1 would be considerably harder against a black-box combination of many weak labelers) but removes Snorkel-style noise-cancellation-via-redundancy as a mitigation — a tradeoff made explicit rather than left implicit, and the direct motivation for Section 4.3’s scope note running through every downstream section.

**Conformal prediction and out-of-distribution detection.** Split conformal prediction [Vovk et al., 2005, Lei et al., 2018] and its distribution-free coverage guarantee (Section 7) have seen growing NLP and security application as a distribution-free alternative to ad hoc confidence thresholds; Angelopoulos and Bates [2023] survey the modern toolkit this paper’s Section 7.1 draws on directly. Out-of-distribution detection (Section 8) via softmax-confidence baselines [Hendrycks and Gimpel, 2017], energy scores [Liu et al., 2020], and class-conditional embedding distance [Lee et al., 2018] is comparatively less explored specifically for weak-label security-NLP classifiers, where – as Section 8.2 shows concretely – the usual confidence-based methods can fail in a domain-specific way (a classifier’s most confident predictions coinciding with exactly the “nothing detected” signature that off-topic input also produces).

### 3 Dataset Construction

RansKnow [Kabuye, 2026] consists of 1,128 video records sourced from 92 cybersecurity YouTube channels, of which 1,034 have an extracted transcript with real content; the remaining 94<sup>1</sup> could not be transcribed by any method in the fetch cascade.

#### 3.1 Channel selection

Channels were curated in two batches. Batch 1 (50 channels, C01–C50) covers established DFIR vendors, security conferences, and independent researchers. Batch 2 (51 candidate channels, C51–C101) expanded coverage into threat-intelligence vendors, additional conferences, and OT/ICS-security and government channels; 42 of the 51 candidates yielded usable transcripts, 1 was a confirmed duplicate of an existing Batch 1 channel, 4 had no discoverable dedicated YouTube channel, and 4 resolved to real channels with zero ransomware-keyword-matching content in their upload history. Full channel-level provenance for both batches is in rubrics/Dataset\_Channel\_Registry\_Updated\_50\_fixed\_urls.xlsx (Batch 1) and rubrics/Dataset\_Channel\_Registry\_Batch2.xlsx (Batch 2), released alongside the corpus rather than reproduced inline as a 92-row table.

#### 3.2 Video selection

Within each channel, videos were selected by scanning recent uploads for a fixed 21-term keyword list (ransomware, extortion, lockbit, incident response, ...; full list in Appendix A) matched against title and description, capped at 20 videos per channel and a minimum duration of 5 minutes (short clips and trailers excluded). This is a convenience sample, not a random or exhaustive one, and the selection mechanism itself introduces bias worth naming rather than assuming away: a video is only in scope if its *title or description* happens to contain one of the 20 terms, which conditions the corpus on each channel operator’s own titling and SEO/description-writing habits. A channel that discusses ransomware extensively but titles episodes generically (e.g. by client-engagement number rather than threat name) would be under-represented regardless of its actual content relevance, and channels that lean on a small number of high-profile named families (lockbit, conti, ...) in titles for engagement will be over-represented relative to channels discussing less headline-grabbing groups. This compounds with Section 10’s broader sampling-bias point: keyword-based inclusion is a second, independent bias layer on top of channel curation, not a neutral technical filtering step.

<sup>1</sup>Verified directly against the released transcripts/ tree, not assumed from 1128 – 1034: all 94 are confirmed failed transcript extraction (no method in the fetch cascade, Section 3.3, produced usable text), not videos excluded post-fetch for another reason.

### 3.3 Transcript acquisition

Transcripts were acquired via a three-stage cascade: (1) the YouTube Data API’s official captions where available; (2) yt-dlp’s auto-generated captions; (3) local Whisper transcription as a fallback. **This cascade matters more than a footnote:** only 74 of 1,034 transcripts (7.2%) came from official captions. 959 (92.7%) are ASR output – 554 from whisper\_base (the smallest, least accurate Whisper checkpoint) and 405 from a local Whisper fallback, 1 from yt-dlp captions. This raw composition is Finding B of the label-quality audit (Section 4); the downstream implication – what an ASR-dominated corpus does to jargon-term detection specifically – is reported there rather than duplicated here.

## 4 Label-Quality Audit

The Knowledge Agent pipeline (Scripts/knowledge\_agent.py) extracts structured features from every transcript via regex matching: ransomware family mentions (against a maintained alias list), MITRE ATT&CK tactic mention counts across 10 tactics, offensive tool mentions across 8 tools, and platform signals (Windows/Linux/ESXi). Three concrete problems were found and are described here rather than left as an implicit caveat, because a downstream classifier trained on these labels inherits whatever bias they contain.

### 4.1 Finding A: generic-word alias collisions

Nine of 41 family aliases collide with common English or business words (play, hive, cactus, agenda, clop, inc, maze, medusa, royal) – identified programmatically via dictionary lookup against /usr/share/dict/words, then confirmed by manual transcript inspection (e.g. hive matching “registry hive,” a routine DFIR term, rather than Hive ransomware). A windowed-context disambiguation fix (requiring a ransomware-context term within ~15 tokens of the match) was applied and reduces false-positive volume substantially without eliminating it – residual false positives still occur when context vocabulary is itself dense (see Scripts/audit\_family\_aliases.py for the audit methodology and Table 1 for before/after counts).

Table 1: Before/after family mention counts for the 9 flagged aliases, windowed-context disambiguation applied. Play alone accounted for 386 of 494 (78%) of all family-tagged videos before the fix.

| Family | Before | After | Dropped |
|--------|--------|-------|---------|
| Play   | 386    | 111   | 275     |
| Qilin  | 88     | 21    | 67      |
| Hive   | 23     | 10    | 13      |
| Maze   | 11     | 5     | 6       |
| Royal  | 28     | 19    | 9       |
| INC    | 10     | 7     | 3       |
| Cactus | 9      | 6     | 3       |
| Cl0p   | 11     | 11    | 0       |
| Medusa | 7      | 7     | 0       |

### 4.2 Finding B: ASR-dominated transcript quality

As noted in Section 3.3, 92.7% of transcripts are ASR output rather than official captions.

**[TODO]** Run the Phase 0.4 spot-check (manual comparison of ASR output against source audio for a fixed jargon term list – LockBit, Mimikatz, PsExec, Cobalt Strike, rclone, BloodHound – across the whisper-transcribed subset) and report a measured word-error-rate-on-jargon figure here. Aggregate detection-rate comparisons across Transcript\_Provider did not show an obvious gap, which is not the same as no effect – individual mistranscriptions of rare jargon can matter without moving an aggregate rate.

### 4.3 Scope: no independent ground truth

An earlier design stage of this project planned a held-out, independently annotated evaluation subset (180 videos, stratified by channel type, year, and current family-count) specifically to break the circularity of evaluating a classifier against the same weak-label pipeline it was trained on. That subset was never annotated. External annotator recruitment required institutional ethics approval not available to this project; a single-investigator pass was considered as a fallback and rejected as still too close to the system under evaluation to serve as independent ground truth in a way worth claiming in a paper.

The consequence is stated here plainly, and repeated at the point every affected result is reported, rather than left as a single Limitations bullet a reader could miss: **every quantitative result in Section 5 onward, with the explicit exception of Section 8’s out-of-distribution results, measures agreement with the audited Knowledge Agent extraction pipeline (Section 4), not verified accuracy against human judgment.** Section 4 demonstrates concretely that this pipeline’s labels contain real, non-trivial error even after the disambiguation fix (Finding A); nothing downstream corrects for that. Readers should treat macro-F1 and calibration numbers in this paper as characterizing how well models reproduce a specific, imperfect, documented extraction process – not as claims about ransomware-domain understanding in an absolute sense. Section 10 returns to this as the paper’s central limitation.

## 5 Task Definitions and Baselines

Five prediction tasks are defined over the corpus: ransomware family (multi-label, 18 classes after excluding 3 singleton-occurrence classes, see Table 2), dominant MITRE tactic (multi-class, 11 classes), platform signal (multi-class, 4 classes), offensive tool mentions (multi-label, 7 classes after excluding MegaNZ, which has zero occurrences in the corpus), and a ransomware-relevance binary screen (derived, not a stored column: `Family_Count > 0 OR Tactic_Impact > 0`).

Table 2: Task taxonomy.

| Task                | Type                | Classes | Label source                      |
|---------------------|---------------------|---------|-----------------------------------|
| Family              | multi-label         | 18      | alias regex match (audited, §4.1) |
| Dominant tactic     | multi-class (+none) | 11      | max of 10 tactic-pattern counts   |
| Platform signal     | multi-class (+none) | 4       | max of 3 platform-pattern counts  |
| Tool mentions       | multi-label         | 7       | per-tool regex match              |
| Ransomware-relevant | binary              | 2       | derived                           |

### 5.1 Model ladder and validation protocol

Each task is evaluated over three feature representations (structured Knowledge Agent columns, TF-IDF, mean-pooled sentence embeddings<sup>2</sup>) and five models (the rule-based Knowledge Agent itself as baseline 0, Logistic Regression, Linear SVM, Random Forest, LightGBM), under three cross-validation protocols: stratified random  $k$ -fold, GroupKFold by `Channel_ID` (to detect channel-style leakage – a model can learn “this is CrowdStrike’s narration style” instead of the content), and a temporal split (train pre-2024, test 2024+, since 79% of the corpus is from 2024 onward).

### 5.2 Results

Per Section 4.3, every number in this section is weak-label agreement – how well a model reproduces the Knowledge Agent’s own extraction, not verified accuracy. The rule-based Knowledge Agent baseline scores a tautological macro-F1 of 1.000 on every task (checked against its own labels by construction; included as baseline 0, not as evidence of anything) and is omitted from Table 3 accordingly.

<sup>2</sup>384-dim all-MiniLM-L6-v2, chunked at ~200 words per chunk and mean-pooled, computed via a dedicated Python environment (torch 2.5.1) since the reference implementation’s default environment ships an older torch that cannot satisfy `sentence_transformers`’ requirements; see `Scripts/precompute_embeddings.py`.

Table 3: Best stratified-CV macro-F1 per task, across every representation (structured / TF-IDF / embedding) and trained model (LogReg, Linear SVM, Random Forest, GBM) tried. Weak-label agreement, not verified accuracy (§4.3).

| Task                | Best representation | Best model | Macro-F1 |
|---------------------|---------------------|------------|----------|
| Ransomware-relevant | TF-IDF              | GBM        | 0.903    |
| Dominant tactic     | structured          | GBM        | 0.813    |
| Platform signal     | TF-IDF              | GBM        | 0.606    |
| Tool mentions       | TF-IDF              | LogReg     | 0.200    |
| Family              | embedding           | LogReg     | 0.129    |

Two patterns worth naming, both consistent with Section 9’s independent checks: (1) tree models (GBM in particular) are the strongest model wherever tested, and structured features remain the strongest representation specifically for `dominant_tactic` – unsurprising, since those features are numerically close to the counts the label itself was derived from; (2) `family` and `tool` are weak under every representation and model combination tried (`family` never exceeds 0.13 macro-F1; `tool` is additionally unstable across CV splits, ranging 0.03–0.20). This is not obviously a framework or tuning limitation – Section 9’s length and channel-identity confound checks rule out two of the more common alternative explanations for weak tabular/text-classification performance, leaving weak-label ceiling (§4.3) and genuine class sparsity (Section 5’s 18-class `family` / 7-class `tool` taxonomies, several classes with single-digit support) as the remaining candidate explanations.

Table 3 reports only the single best (representation, model) cell per task; the full sweep – 5 tasks  $\times$  3 representations  $\times$  4 trained models  $\times$  3 CV protocols, 180 combinations in total (outputs/experiments/phase1\_\*.csv) – is summarized visually in Figure 1, collapsing the representation axis by taking the best of the three for each (task, model, split) cell so the figure stays legible, and given in full, exact numbers for the stratified-random split in Table 4.




Figure 1: Best-representation macro-F1 for every (task, model) pair, one panel per CV protocol (weak-label agreement, §4.3). `family/tool` stay weak under every model and every split; `relevance` is both the strongest task and the most split-robust.

Table 4: Best-representation macro-F1 per (task, model), stratified-random split (weak-label agreement, §4.3). Channel-grouped and temporal splits follow the same pattern with a modest across-the-board drop, per Figure 1.

| Task            | LogReg | SVM  | RF   | GBM  |
|-----------------|--------|------|------|------|
| Relevance       | 0.71   | 0.68 | 0.82 | 0.90 |
| Dominant tactic | 0.74   | 0.80 | 0.67 | 0.81 |
| Platform        | 0.41   | 0.39 | 0.40 | 0.61 |
| Tool            | 0.20   | 0.18 | 0.04 | 0.07 |
| Family          | 0.13   | 0.08 | 0.02 | 0.08 |

Figure 2 plots Table 4 as a 3D stem plot rather than solid bars: each (task, model) cell is a thin marked stem rather than an opaque block, so no value is hidden behind a taller neighbor from any viewing angle, which a solid-bar rendering of the same data could not guarantee.

Baseline sweep in 3D: best-representation macro-F1 by task and model  
(stratified-random split, weak-label agreement, Section 4.3)




Figure 2: Table 4 as a 3D stem plot – every point visible regardless of viewing angle, unlike a solid-bar rendering of the same data.

## 6 Uncertainty Quantification

**Scope note (§4.3 applies here directly):** calibration is normally reported against ground truth – does an 80%-confidence prediction turn out correct roughly 80% of the time. Without independent labels, the only calibration target available here is the weak-label distribution itself, which changes what a well-calibrated result would mean: a model perfectly calibrated against its own training signal has learned to be a confident, well-calibrated echo of a regex, not necessarily a well-calibrated model of ransomware content. This section is retained, with that meaning made explicit at every reported number, because reliability diagrams and Expected Calibration Error against the weak labels are still informative about a narrower question worth answering: does a model’s confidence track its *agreement with the extraction pipeline*, or is it uniformly overconfident regardless of agreement. The latter would be a red flag on its own.

### 6.1 Calibration against weak labels

For a multiclass task, let  $\hat{p}(\cdot | x)$  be the trained model’s predicted class distribution for input  $x$ , and let  $\hat{c}(x) = \arg \max_c \hat{p}(c | x)$  with confidence  $\hat{p}_{\max}(x) = \max_c \hat{p}(c | x)$ . Agreement with the weak label  $y^{\text{KA}}(x)$  (the Knowledge Agent’s own extraction) is  $a(x) = \mathbb{I}[\hat{c}(x) = y^{\text{KA}}(x)]$ . For the two multi-label tasks (family, tool), whose models are independent per-class binary classifiers (Section 5), each (sample, class) pair  $(x, k)$  is treated as its own calibration point, with confidence  $\hat{p}_{\max}(x, k) = \max(\hat{p}_k(x), 1 - \hat{p}_k(x))$  and agreement against the binary weak label thresholded at 0.5.

Reliability is summarized with  $B=10$  equal-width confidence bins  $\{I_b\}_{b=1}^B$  and the standard Expected Calibration Error,

$$\text{ECE} = \sum_{b=1}^B \frac{|S_b|}{N} |\text{acc}(S_b) - \text{conf}(S_b)|, \quad S_b = \{x : \hat{p}_{\max}(x) \in I_b\}, \quad (1)$$



where  $\text{acc}(S_b) = \frac{1}{|S_b|} \sum_{x \in S_b} a(x)$  and  $\text{conf}(S_b) = \frac{1}{|S_b|} \sum_{x \in S_b} \hat{p}_{\max}(x)$ . Per the scope note above,  $a(x)$  measures agreement with  $y^{\text{KA}}$ , not correctness, so Eq. (1) is a *weak-label ECE*. The headline number below the scope note is the confident-disagreement rate at threshold  $\tau = 0.9$ ,

$$\text{CDR}_{\tau} = \frac{1}{N} \sum_{i=1}^N \mathbb{I}[\hat{p}_{\max}(x_i) > \tau] \cdot (1 - a(x_i)), \quad (2)$$

the closest available analogue to a “confidently wrong” rate without independently verified labels to actually be wrong against.

Each task is evaluated on the 5-fold stratified out-of-fold predictions of its best (feature, model) combination from Table 3, except family, whose best combination (embedding representation) needs the dedicated environment described in Section 5’s footnote and is reported without a calibration curve here.

Table 5: Weak-label calibration (Equations (1) and (2)) and the RF tree-vote entropy proxy (Equation (3)), best (feature, model) per task.  $N$  is the number of (sample) or (sample, class) calibration points. Entropy AUROC is how well entropy alone separates agreement from disagreement with the weak label (Algorithm 1).

| Task               | $N$  | ECE   | $\text{CDR}_{0.9}$ | Entropy (agree/disagree) | Entropy AUROC |
|--------------------|------|-------|--------------------|--------------------------|---------------|
| Relevance          | 1034 | 0.079 | 0.071              | 0.910 / 0.980            | 0.769         |
| Dominant tactic    | 1034 | 0.051 | 0.043              | 0.275 / 0.551            | 0.860         |
| Platform           | 1034 | 0.110 | 0.104              | 0.717 / 0.769            | 0.675         |
| Tool (per-label)   | 7238 | 0.016 | 0.016              | 0.121 / 0.271            | 0.977         |
| Family (per-label) | 1034 | –     | –                  | 0.095 / –                | –             |

Figure 3 plots the reliability curves underlying Table 5. Two things worth naming honestly rather than smoothing over: first, per-bin counts are small (single digits to low tens per bin at  $N=1034$ ), so individual bin heights are noisy – the curves are directionally informative, not point-estimate-precise. Second, family’s row has no ECE/CDR: its calibration point is unavailable in this environment (as above), so only the RF tree-vote entropy proxy, computed on the TF-IDF fallback representation, is reported for it, and even that proxy has no “disagree” entries to average – at the per-(sample, class) level, most of family’s 18 classes are negative for most samples, so the majority-class prediction (“absent”) agrees with the weak label on essentially every RF-on-TF-IDF calibration point despite the task’s low macro-F1 in Table 3, which weights all 18 classes equally rather than by prevalence. Both readings are correct simultaneously; they are just answering different questions (per-pair agreement vs. per-class balanced performance).




Figure 3: Reliability diagrams against weak labels, one panel per task with a usable calibration curve (Table 5). Diagonal is perfect calibration.

## 6.2 A free uncertainty proxy: Random Forest tree-vote entropy

Reliability diagrams need a held-out calibration protocol; a Random Forest offers a per-prediction uncertainty signal with no separate calibration step. For an input  $x$  and a forest of  $T$  trees producing



individual (hard) predictions  $t_1(x), \dots, t_T(x)$  over  $K$  classes, let  $\pi_c(x) = \frac{1}{T} \sum_{j=1}^T \mathbb{1}[t_j(x) = c]$  be the empirical vote share for class  $c$ . The normalized tree-vote entropy is

$$H(x) = -\frac{1}{\log K} \sum_{c=1}^K \pi_c(x) \log \pi_c(x) \in [0, 1], \quad (3)$$

with  $H(x)=0$  when every tree agrees and  $H(x)=1$  at a uniform split across all  $K$  classes. For the multi-label tasks,  $K=2$  per class and  $H(x, k)$  is averaged over classes  $k$  for a per-sample summary. Algorithm 1 is the exact out-of-fold procedure used to produce Table 5’s entropy columns.

---

**Algorithm 1** RF tree-vote entropy, 5-fold out-of-fold

---

**Require:** dataset  $\mathcal{D}$ , task with label function  $y(\cdot)$ , feature builder  $\phi$

- 1:  $\mathcal{F} \leftarrow 5$  stratified folds of  $\mathcal{D}$
- 2: **for** fold  $(\mathcal{D}_{\text{train}}, \mathcal{D}_{\text{test}}) \in \mathcal{F}$  **do**
- 3:   fit  $\phi$  on  $\mathcal{D}_{\text{train}}$ ; fit Random Forest  $\{t_j\}_{j=1}^T$  on  $\phi(\mathcal{D}_{\text{train}})$
- 4:   **for**  $x \in \mathcal{D}_{\text{test}}$  **do**
- 5:      $t_j(x) \leftarrow$  individual tree prediction,  $j = 1, \dots, T$   $\triangleright$  sklearn returns internal integer-encoded classes here, not the original label strings – see note below
- 6:     recover original class labels via the fitted forest’s class index before comparing to  $y(x)$
- 7:     compute  $H(x)$  via Equation (3); record  $a(x) = \mathbb{1}[\text{forest majority vote} = y(x)]$
- 8:   **end for**
- 9: **end for**
- 10: **return**  $\{(H(x), a(x))\}_{x \in \mathcal{D}}$ , pooled across folds

---

The in-line comment in Algorithm 1 records a real implementation pitfall, not a hypothetical one: scikit-learn fits each tree in a `RandomForestClassifier` on internally label-encoded integer targets, so `tree.predict()` returns class *indices*, not the original label strings the ensemble’s own `.predict()` returns. Comparing tree-level votes directly against string labels without mapping back through the fitted forest’s class index silently produces 0% apparent agreement on every multiclass task – caught in this project’s own Phase 2 run (initial numbers showed exactly that pattern across all three multiclass tasks simultaneously, which is what made it legible as a bug rather than a real result) and is recorded here as a concrete instance of the caution Section 10 urges about verifying rather than assuming standard library internals. Table 5’s entropy AUROC column shows the proxy works well once fixed: 0.68–0.98 across tasks, best on `tool` (many easy true negatives) and weakest on `platform`.

## 7 Conformal Prediction

**Scope note:** split conformal prediction’s distribution-free coverage guarantee is a statement about exchangeability with the calibration set, not about label correctness – it holds formally even when calibrated against weak labels, because the guarantee is “the true *weak* label falls in the predicted set  $(1 - \alpha)$  of the time,” not “the true *real-world* label does.” That distinction is the entire content of this scope note: conformal coverage against the Knowledge Agent’s labels is a well-defined, honestly reportable number, but it validates agreement with the extraction pipeline, not correctness. Reported as such.

### 7.1 Method: split conformal with the LAC score

Split conformal prediction (Vovk et al., 2005, Lei et al., 2018; see Angelopoulos and Bates, 2023 for a tutorial treatment) uses a nonconformity score  $s(x, y)$  that is small when  $y$  is a plausible label for  $x$ . This paper uses the least-ambiguous-set (LAC) score,

$$s(x, y) = 1 - \hat{p}(y \mid x), \quad (4)$$

calibrated on a held-out split  $\mathcal{D}_{\text{calib}}$  of size  $n$  disjoint from both the training and test data. The calibration threshold uses the standard finite-sample correction,

$$\hat{q} = \text{Quantile}\left(\{s(x_i, y_i)\}_{i=1}^n; \frac{\lceil (n+1)(1-\alpha) \rceil}{n}\right), \quad (5)$$

and the prediction set for a new input  $x$  is

$$C(x) = \{y : \hat{p}(y | x) \geq 1 - \hat{q}\}. \quad (6)$$

Under exchangeability of  $\mathcal{D}_{\text{calib}}$  and the test point, this construction guarantees

$$\mathbb{P}[y_{\text{test}} \in C(x_{\text{test}})] \geq 1 - \alpha, \quad (7)$$

with  $\alpha=0.10$  (90% target coverage) throughout. As the scope note above states,  $y_{\text{test}}$  here is the weak label  $y^{\text{KA}}$ : Equation (7) is a real, distribution-free guarantee about agreement with the extraction pipeline, not about real-world correctness. For `dominant_tactic` and `platform`, Equations (4) to (6) apply directly over the task’s full class vocabulary. `family` and `tool` have no single multiclass structure to conformalize – each is an independent per-class binary classifier (Section 5) – so each class  $k$  gets its own binary LAC calibration and its own  $\hat{q}_k$ , and a sample’s per-label prediction sets are reported both individually and jointly (all labels simultaneously covered).

Algorithm 2 is run for 5 repeated random 60/20/20 train/calibration/test splits per task, results averaged across repeats rather than reported from one arbitrary split – a single 20% calibration set at this sample size ( $n \approx 200$ ) is not a reliable enough estimate of  $\hat{q}$  to report alone.

---

**Algorithm 2** Split conformal calibration and coverage evaluation

---

**Require:** dataset  $\mathcal{D}$ , task, target coverage  $1 - \alpha$ , repeats  $R$

- 1: **for**  $r = 1, \dots, R$  **do**
  - 2:   split  $\mathcal{D}$  into  $\mathcal{D}_{\text{train}}$  (60%),  $\mathcal{D}_{\text{calib}}$  (20%),  $\mathcal{D}_{\text{test}}$  (20%), stratified
  - 3:   fit model on  $\mathcal{D}_{\text{train}}$
  - 4:   compute  $\hat{q}$  on  $\mathcal{D}_{\text{calib}}$  via Equations (4) and (5)
  - 5:   **for**  $x \in \mathcal{D}_{\text{test}}$  **do**
  - 6:     build  $C(x)$  via Equation (6); record  $\mathbb{I}[y^{\text{KA}}(x) \in C(x)]$  and  $|C(x)|$
  - 7:   **end for**
  - 8: **end for**
  - 9: pool test-set coverage indicators across all  $R$  repeats
  - 10: report marginal coverage, and coverage restricted to the  $\text{Year} \geq 2026$  and Conference-channel slices
- 

## 7.2 Results

Table 6: Split-conformal coverage at 90% target ( $\alpha=0.10$ ), pooled over 5 repeats,  $n \approx 1035$  pooled test predictions per task. `family/tool` report mean per-label coverage and the stricter all-labels-jointly-covered rate. Set size is out of the task’s class count ( $K$ ) for `dominant_tactic/platform`, or out of 2 (per label) for `family/tool`.

| Task               | $K$ | Coverage (marginal)  | Avg. set size | Slice: Yr $\geq 26$    | Slice: Conf.           |
|--------------------|-----|----------------------|---------------|------------------------|------------------------|
| Dominant tactic    | 11  | 0.877                | 0.94/11       | 0.890                  | 0.870                  |
| Platform           | 4   | 0.925                | 1.19/4        | 0.932                  | 0.899                  |
| Family (per-label) | 18  | 0.912 (joint: 0.446) | 0.93/2        | 0.487 <sup>joint</sup> | 0.572 <sup>joint</sup> |
| Tool (per-label)   | 7   | 0.909 (joint: 0.650) | 0.92/2        | 0.640 <sup>joint</sup> | 0.752 <sup>joint</sup> |

Figure 4 plots Table 6 directly. Reading the results against the two questions this section set out to answer: first, does the guarantee hold empirically – `dominant_tactic` and `platform`’s marginal coverage (0.877, 0.925) both sit close to the 0.90 target, with set sizes well under one class on average (0.94/11, 1.19/4), i.e. the conformal sets are usually a single class or empty, not the whole vocabulary hedged defensively. Second, the conditional-coverage stress test: neither the  $\text{Year} \geq 2026$  nor the Conference-channel slice shows the kind of collapse that would flag a real distribution-shift problem for these two tasks (all four slice numbers stay within about 3 points of the marginal rate). `family` and `tool`’s per-label coverage (0.91, 0.91) is equally close to target, but their *joint* (all-labels-simultaneously) coverage is far lower (0.45, 0.65) – an expected consequence of compounding per-label error rates across 18 and 7 largely-independent binary calibrations respectively, not a failure of Equation (7) (which was never a joint guarantee to begin with) and not itself evidence of conditional-coverage collapse, since the joint-coverage slices track the joint-coverage marginal

about as closely as the multiclass tasks’ slices track theirs. Per Section 7.1, all of the above describes agreement with  $y^{KA}$ , and Section 8 below is the complementary analysis that does not carry that caveat.




Figure 4: Empirical conformal coverage, marginal vs. the two conditional-coverage stress-test slices, against the 90% target (Table 6).

## 8 Out-of-Distribution Detection

Unlike Sections 5–7, this section’s results are not subject to the weak-label scope note in Section 4.3: OOD detection asks whether a model’s own confidence or embedding-space behavior differs between in-distribution and out-of-distribution input, which is answerable without knowing whether either input’s *label* is correct.

### 8.1 Setups and methods

Four setups, each with a known-by-construction in-distribution (ID)/out-of-distribution (OOD) split, drawn from real structure in the corpus rather than synthetic corruption:

1. **Temporal emergence.** ID: pre-2024 videos. OOD: post-2024 videos mentioning a family whose first appearance in this corpus is 2024 or later – verified directly against per-family first-occurrence year rather than assumed (Maze, Medusa, Black Basta, INC: 2024; Cactus: 2025), at Algorithm 3’s assertion step.
2. **Channel-type shift.** ID: Vendor + DFIR channels. OOD: Conference + Independent channels.
3. **Leave-rare-family-out.** ID: videos never mentioning “Cl0p” (11 occurrences total – long-tailed, unlike Conti/LockBit/ Qilin/Akira, each with 19+). OOD: the 11 Cl0p-mentioning videos.
4. **External corpus.** ID: RansKnow transcripts. OOD: 20 Newsgroups (rec.sport.baseball + sci.space), included as a sanity check that the methods below can detect an obviously unrelated domain before trusting them on setups 1–3.

Four detection methods, computed from a shared `dominant_tactic` classifier (structured features, LightGBM – Table 3’s strongest non-circular result, and a natural multiclass fit for the softmax-family methods below), refit fresh on each setup’s ID-only training pool so that no OOD signal leaks into

fitting:

$$s_{\text{MSP}}(x) = 1 - \max_c \hat{p}(c | x) \quad (8)$$

$$s_{\text{energy}}(x) = -\log \sum_c \exp(f_c(x)) \quad (9)$$

$$s_{\text{maha}}(x) = \min_c (\phi(x) - \mu_c)^\top \Sigma^{-1} (\phi(x) - \mu_c) \quad (10)$$

$$s_{\text{conf}}(x) = 1 - \hat{p}(\hat{c}(x) | x) \quad (\text{Equation (4) at the predicted class}) \quad (11)$$

Equation (8) is Hendrycks and Gimpel’s max-softmax-probability baseline [Hendrycks and Gimpel, 2017]. Equation (9) is the energy score [Liu et al., 2020], computed from LightGBM’s raw pre-softmax margins  $f_c(x)$  (not from  $\hat{p}(c | x)$ : probabilities already sum to 1, which makes Equation (9) degenerate to a constant if computed from them directly – a real pitfall caught while implementing this section, not a hypothetical one). Equation (10) is the class-conditional Gaussian Mahalanobis distance of Lee et al. [2018], with  $\phi(x)$  the 384-dim mean-pooled sentence embedding (Section 5’s footnote) and  $\mu_c, \Sigma$  the ID-train class-conditional means and pooled covariance, using the classifier’s own *predicted* classes to form groups (consistent with Section 4.3: no independently verified labels exist to group by instead). Equation (11) reuses Section 7’s LAC score, evaluated at the predicted rather than true class for OOD flagging – this is deliberately, not accidentally, the same quantity as Equation (8); the two are reported separately for comparability with Equation (5)’s calibrated threshold (does a point’s own score exceed the 90%-target  $\hat{q}$ ), which is the operationally distinct number.

---

#### Algorithm 3 OOD setup evaluation

---

**Require:** corpus  $\mathcal{D}$ , ID/OOD partition rule for the setup

- 1: **assert** every claimed-emerging family’s first-occurrence year matches the setup’s stated cutoff  $\triangleright$  caught a factual error at this step – see note below
  - 2: split ID pool into  $\mathcal{D}_{\text{train}}$  (75%) /  $\mathcal{D}_{\text{test}}^{\text{ID}}$  (25%)
  - 3: fit `dominant_tactic/structured/LightGBM` on  $\mathcal{D}_{\text{train}}$
  - 4: score  $\mathcal{D}_{\text{test}}^{\text{ID}}$  and  $\mathcal{D}_{\text{OOD}}$  under Equations (8) to (11)
  - 5: report AUROC(ID vs. OOD) per method, plus the conformal exceedance-rate contrast
- 

The assertion in Algorithm 3 is not decorative: the temporal- emergence setup’s family list was originally drafted from memory as “Qilin, DragonForce, Black Basta,” before checking actual first-occurrence years against the corpus. Qilin’s first real occurrence is 2019, not post-2023, and DragonForce does not appear in this corpus at all – both would have silently produced a meaningless setup (there is no real ID/OOD separation if the “emerging” families already appear throughout the ID training period). The corrected list (Maze, Medusa, Black Basta, INC, Cactus) is the one actually used in Table 7 and is verified by Algorithm 3 on every run, not just checked once by hand.

## 8.2 Results

Table 7: OOD-detection AUROC by setup and method (Equations (8) to (11)).  $n_{\text{ID}}/n_{\text{OOD}}$  are ID-test/OOD sample counts. Mahalanobis is unavailable for `external_corpus` (20 Newsgroups has no entry in the RansKnow-only embedding cache).

| Setup                    | $n_{\text{ID}}$ | $n_{\text{OOD}}$ | MSP   | Energy | Mahalanobis | Conformal |
|--------------------------|-----------------|------------------|-------|--------|-------------|-----------|
| Temporal emergence       | 55              | 31               | 0.561 | 0.360  | 0.223       | 0.533     |
| Channel-type shift       | 156             | 353              | 0.477 | 0.487  | 0.666       | 0.487     |
| Leave-rare-family-out    | 256             | 11               | 0.377 | 0.350  | 0.145       | 0.376     |
| External corpus (sanity) | 259             | 1190             | 0.625 | 0.476  | –           | 0.631     |

Figure 5 plots Table 7. The honest reading is a mixed and mostly negative one, which is itself the finding worth reporting rather than a result to explain away. Three observations:

**The confidence-based methods (MSP, energy, conformal) hover near chance (0.5) on every setup, including the sanity check.** Diagnosing *why* on `external_corpus` is informative: 99.6% of

20 Newsgroups documents are classified None (no dominant tactic detected) by the shared classifier, and with *higher* average confidence than the real in-distribution test set (mean  $\hat{p}_{\max}$ : 0.921 OOD vs. 0.895 ID). A near-empty, off-topic TF-IDF vector and a genuinely tactic-free ransomware transcript both produce the same signature – confidently “no tactics detected” – so a classifier whose decision function is fundamentally an argmax over small keyword-mention counts has no way to distinguish “nothing here” from “wrong domain entirely.” This is a known blind spot of softmax/probability-based OOD detection in general [Hendrycks and Gimpel, 2017], reproduced concretely here rather than assumed from the literature.

**Mahalanobis distance on real sentence embeddings is the one method that separates classes at all, and only partially.** It reaches 0.666 on channel-type shift – Conference/Independent channels differ enough in embedding space (production style, framing, audience) to be mildly detectable – but performs *worse* than chance on temporal emergence (0.223) and leave-rare-family-out (0.145). Both of those setups’ OOD population describes the same underlying phenomenon (ransomware incidents, attacker tradecraft) as the ID population, just naming a different or newer group – exactly the case where sentence embeddings, which encode semantic content rather than proper nouns, would *not* be expected to separate the two populations well. The small OOD sample sizes in these two setups ( $n=31$ ,  $n=11$ ) also mean each AUROC estimate carries substantial variance on its own; with  $n_{\text{OOD}}=11$ , a single re-ordered pair moves the estimate by roughly  $1/11$ .

**The conformal exceedance-rate contrast (Equation (5)’s  $\hat{q}$ ) is uninformative in this corpus specifically.** In every setup, the calibrated threshold  $\hat{q}$  lands so close to the maximum possible score ( $\hat{q} \geq 0.99$  in three of four setups) that essentially no point, ID or OOD, exceeds it – a consequence of the classifier being very confidently correct on the bulk of its small ( $n \approx 115$ ) calibration split, which pushes the 90th-percentile score to an extreme value. Equation (11)’s AUROC column is the more informative number from this method here; the exceedance-rate contrast is reported in the released JSON (outputs/phase4\_ood.json) for completeness but is not a useful discriminator at this sample size.




Figure 5: OOD-detection AUROC by method (color) and setup (axis), against the chance line at 0.5 (Table 7). Mahalanobis is the only method to clear chance by a wide margin, and only on one of four setups.

A single AUROC number per (method, setup) cell is a summary of a comparison between two full score *distributions*, and the two figures below show those distributions directly rather than only their summary statistic. Figure 6 plots the actual MSP score density for the ID-test and OOD populations, per setup, as a 3D ridgeline: the near-total overlap between the blue (ID) and red (OOD) curves in every setup is the direct visual counterpart of the near-chance AUROC numbers above, and makes concrete what “the classifier’s most confident predictions coincide with the off-topic-input signature” actually looks like as two overlapping peaks near the same low-score, high-confidence region, not two separated ones. Figure 7 shows the complementary picture in representation space rather than

score space: a 3D t-SNE projection of the same 384-dim sentence embeddings the Mahalanobis method (Equation (10)) computes its distances on, colored by ID/OOD membership. OOD points are visibly scattered throughout the ID cloud with no separable structure in `temporal_emergence` and `leave_rare_family_out` – consistent with those setups’ below-chance Mahalanobis AUROC (0.223, 0.145) – while `channel_type_shift` shows OOD points modestly concentrated toward one region of the embedding cloud, consistent with that setup being the one case where Mahalanobis distance actually clears chance (0.666).




Figure 6: ID-test vs. OOD max-softmax-probability score density per setup, as a 3D ridgeline (kernel density estimate, not binned histograms). Overlap between the blue and red curves is the direct visual counterpart of near-chance AUROC.




Figure 7: 3D t-SNE of the 384-dim sentence embeddings per setup, colored by ID/OOD membership (`external_corpus` omitted – no cached embedding for 20 Newsgroups). No visible separation in the two setups with near-chance Mahalanobis AUROC; modest clustering in channel-type shift, the one setup where it clears chance.



Taken together, Section 8.2 does not support a general claim that this pipeline’s models reliably detect distributional shift; it supports a narrower, still useful one: OOD detection here is hard, softmax-confidence-based methods have an identifiable and reproducible blind spot on this specific classifier (the “no signal” vs. “wrong domain” confusion above), and embedding-distance methods help only when the shift is stylistic rather than topical. Unlike Sections 5 to 7, none of this depends on whether  $y^{KA}$  is correct (Section 4.3) – it is a direct, if modest, empirical finding about model behavior.

## 9 Interpretability and Perturbation Robustness

Unlike Sections 5–7, this section asks questions about model *behavior* – does it exploit shortcuts, what does it rely on, how much do its predictions move under semantically-irrelevant edits – none of which require independent ground truth to answer, so Section 4.3’s scope note does not apply here.

### 9.1 Confound checks

**Length.** Transcript word count correlates moderately with `Tactic_Total_Mentions` (Pearson  $r = 0.409$ ) and weakly with `Family_Count` ( $r = 0.131$ ), but not at all with `Tool_Total_Mentions` ( $r = 0.043$ , not significant) – the confound is not uniform across signal types. Because it is mathematically guaranteed that dividing every tactic’s count for one video by that video’s own word count cannot change which tactic has the highest count (argmax is scale-invariant per row), the confound cannot affect how the `Dominant_Tactic` label itself was derived; the question that matters is whether it inflates `dominant_tactic`’s strongest result (structured features, Table 3). Re-scoring with length-normalized (mentions per 1,000 words) instead of raw-count features gives 0.753 macro-F1 against the raw-count model’s 0.764 – a 1.1-point difference, not a length-exploitation artifact.

**Channel identity.** 486 of 1,034 transcripts (47%) contain at least one literal mention of their own channel’s name. Masking those mentions and re-scoring TF-IDF LogReg under both CV protocols leaves the `channel_grouped-vs-stratified_random` gap unchanged for `family` (0.028 raw  $\rightarrow$  0.031 masked) and for `dominant_tactic` (0.020 raw  $\rightarrow$  0.029 masked). Whatever a model exploits when a channel is not held out, it is not the channel literally naming itself – the leakage this split protocol exists to catch is stylistic or topical, not addressable by name-masking.

### 9.2 Counterfactual explanations

Motivated directly by Finding A: does a classifier trained on the *fixed* labels still show the regex’s original word-collision fragility? A TF-IDF+LogReg model was trained on today’s corrected `Family` labels and evaluated on the 275 videos Section 4.1 confirmed as false positives under the pre-fix regex (videos where “play” appears without ransomware context). The trained model assigns essentially zero probability to `Play` on all 275 (maximum  $P = 0.013$ ) – it did not inherit the bare-word-collision bug.

The complementary question – what drives the model’s actual positive `Play` predictions – was then asked with a greedy word-*type* removal counterfactual search (remove the vocabulary word whose removal most reduces  $P(\text{Play})$ , repeat until the prediction flips or a budget is exhausted) on the 8 highest-confidence true positives. `play` appears in 5 of 8 removal sets and is consistently one of the highest-impact single removals when present, but **0 of 8 cases flip on removing “play” alone** – every flip required 4–10 additional removals, most of them generic (function words, not ransomware-specific vocabulary). The model reduced, but did not eliminate, reliance on the single ambiguous trigger word, and the words it combines that reliance with are not obviously a clean, interpretable “ransomware vocabulary” set – worth flagging as its own finding about decision-boundary diffuseness, independent of the word-collision question it was designed to answer.

### 9.3 Perturbation robustness

Untargeted counterpart to the counterfactual search above: two perturbation types, measuring flip rate (does the top prediction change) and mean absolute probability drift on freshly-fit TF-IDF+LogReg models. *Simulated ASR noise* (compound proper nouns split apart, e.g. `WannaCry`  $\rightarrow$  “wanna cry,” `Mimikatz`  $\rightarrow$  “mimi cats,” plus 8% random token deletion) was applied to the 74 official-caption



videos specifically – the “clean” subset, simulating what those transcripts would look like if they had gone through the ASR cascade that produced the other 92.7% of the corpus (Finding B). *Synonym substitution* (common non-jargon words swapped for synonyms) was applied to a 150-video sample.

*dominant\_tactic*, *platform*, and *relevance* are essentially unmoved by either perturbation (0.0–0.7% flip rate, mean drift < 0.01 in every case). *family* flips 20–26% of the time under both perturbations despite a much smaller mean drift (~0.002) than any other task – not a new problem, but the same weak-model issue from Table 3 (*family*’s best macro-F1 is 0.129) surfacing differently: closely-bunched, low-confidence probabilities mean a negligible nudge is enough to change which barely-ahead class wins. On the framing of “adversarial” robustness for this dataset: no real adversary submits transcripts to evade this classifier in production, so a full gradient-based adversarial-attack framework would be disproportionate to what this paper needs; the actual risk surface is ASR transcription noise and channel-style variation, which is what this sweep directly measures. Combined with Section 8’s temporal-emergence and leave-rare-family-out setups, these results are the paper’s answer to how the system handles evolving, novel, and adversarial-adjacent input – not three separate open questions.

## 10 Limitations and Ethics

This paper’s single most consequential limitation is the absence of independent ground truth, discussed at length below; it opens this section as prose, not because the remaining limitations are unimportant, but because they are all of a smaller and more familiar kind (sampling bias, ASR noise, dual use) that most empirical NLP papers carry in some form, whereas the ground-truth gap changes what every quantitative claim in Sections 5–7 *means*, not merely how confidently it can be stated.

- **No independent ground truth.** This is the paper’s central limitation, not one bullet among several, and is why it is stated repeatedly rather than once: this project does not include a human-annotated evaluation set independent of the Knowledge Agent pipeline (Section 4.3). External annotator recruitment required institutional ethics approval not available to this project; a single-investigator annotation pass was considered and rejected as still too close to the system under evaluation to count as independent evidence. Every result in Sections 5–7 is consequently agreement with an audited but demonstrably imperfect extraction pipeline (Finding A), not verified accuracy – Section 8’s OOD results and Section 9’s behavioral diagnostics are the exceptions, since neither depends on label correctness. Recruiting independent annotators, whether through formal ethics approval or an unpaid collaborator arrangement that falls outside whatever definition of human-subjects research applies at the author’s institution, is the single highest-value extension of this work – more valuable than any additional modeling in Sections 5–7 would be, since none of that modeling can currently be checked against anything but itself.
- **Sampling bias.** The corpus is conditioned on which channels operators chose to run, which videos they chose to title and describe in keyword-matchable ways, and YouTube’s own recommendation and search surfacing – not a random or exhaustive sample of ransomware-relevant discourse.
- **ASR noise.** 92.7% of transcripts are machine-transcribed, mostly by *whisper\_base*. Section 9.3’s simulated-ASR-noise sweep bounds how much this moves predictions; Finding B’s word-error-rate-on-jargon question (comparing ASR output against source audio directly, rather than a simulation of ASR-style errors) remains open.
- **Weak-label ceiling.** Every quantitative result in Sections 5–7 is bounded by the Knowledge Agent’s own precision/recall, which Finding A shows is not perfect even after the disambiguation fix, and which Section 9.2 shows partially propagated into at least one trained model.
- **Dual use.** This is a corpus of ransomware tactics, tools, and case studies extracted from public educational content, and it is worth being direct about what that does and does not change. Every underlying video is already public, already indexed by YouTube’s own search, and produced by defenders (vendors, DFIR practitioners, conference speakers) specifically to be watched by other defenders; RansKnow discloses no technique, tool, or tactic that is not already public, and the source material is materially less operationally detailed than existing public resources built for exactly this kind of aggregation – the MITRE ATT&CK knowledge

base itself, and the CTI reporting datasets in Section 2 [Lim et al., 2017, Lange et al., 2024], both already structure comparable tactic/technique/tool information at finer granularity and higher label quality than RansKnow’s weak-label extraction achieves (Section 4.3). What RansKnow adds is not new operational detail but a different modality (transcribed spoken narration rather than written reports) and different scale accessible without specialized CTI tooling. The realistic marginal risk is a small reduction in the effort needed to *aggregate* already-scattered public information into one structured, searchable place, rather than any disclosure of non-public technique detail; this is the same tradeoff every public threat-intelligence aggregation resource makes; and the corpus’s actual utility, given Section 4.3’s weak-label ceiling and Section 5’s weak family/tool classification performance, currently favors defensive research uses (weak-supervision methodology, corpus-level trend analysis) far more than it favors any operational offensive use, for which primary-source technical documentation would be a far more direct and reliable starting point than this corpus’s noisy, ASR-transcribed, regex-extracted labels.

## 11 Conclusion

RansKnow is a 1,128-video, 1,034-transcript, MITRE ATT&CK-aligned corpus of ransomware-related YouTube content, and this paper’s contribution is not the corpus alone. The label-quality audit methodology (Section 4) and the interpretability diagnostics that verify it against a trained model’s actual behavior (Section 9) are a second, reusable contribution: any weak-label security-NLP dataset built the same way (keyword or regex extraction over scraped or transcribed text) can reuse the alias-collision-audit methodology and the counterfactual/ perturbation checks even without reusing this specific corpus, and should – Finding A demonstrates concretely that a plausible-looking regex label can be dominated by false positives (Table 1: 78% of pre-fix Play detections), which is precisely the class of error a corpus released without this kind of audit would carry silently.

The most valuable open follow-up to this specific paper is not further modeling. It is independent, human-annotated verification of a subset of RansKnow’s labels (Section 10), and it is worth being explicit about why that ranks above any modeling extension: every quantitative result this paper reports outside Section 8 would change in *meaning*, not just in number, the day that verification exists. A macro-F1 of 0.813 on `dominant_tactic` currently means “a LightGBM classifier reproduces the Knowledge Agent’s own regex output at this rate”; the same number, computed against even a modest independently-annotated subset, would mean something closer to what a reader of an ML paper’s results table normally assumes a reported score means. Closing that gap changes the paper’s central claim, not just its confidence interval, which is why it is listed first among follow-ups rather than alongside ordinary future-work suggestions like additional model architectures or task definitions.

## Acknowledgments and Disclosure of Funding

This work was supported in part by the Pro:NE (Pronortheast) Network under Grant RDF22/proNE/JHA. The authors thank the members of the CyberNets Research Group, Northumbria University, for their insightful discussions and useful feedback throughout this research project. The opinions, findings, conclusions, or recommendations expressed in this material are solely those of the authors and do not necessarily represent the views of the supporting team. The authors declare no competing interests.

## References

- Anastasios N. Angelopoulos and Stephen Bates. A gentle introduction to conformal prediction and distribution-free uncertainty quantification. *Foundations and Trends in Machine Learning*, 16(4): 494–591, 2023. arXiv:2107.07511.
- Dan Hendrycks and Kevin Gimpel. A baseline for detecting misclassified and out-of-distribution examples in neural networks. In *International Conference on Learning Representations (ICLR)*, 2017.
- Henry Kabuye. RansKnow: A ransomware knowledge dataset. <https://www.kaggle.com/datasets/henrykabuye/ransknow-v1>, 2026. Accessed 2026-08-11.

- Lukas Lange, Marc Müller, Ghazaleh Haratinezhad Torbati, Dragan Milchevski, Patrick Grau, Subhash Chandra Pujari, and Annemarie Friedrich. AnnoCTR: A dataset for detecting and linking entities, tactics, and techniques in cyber threat reports. In *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING 2024)*, 2024.
- Kimin Lee, Kibok Lee, Honglak Lee, and Jinwoo Shin. A simple unified framework for detecting out-of-distribution samples and adversarial attacks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2018.
- Jing Lei, Max G’Sell, Alessandro Rinaldo, Ryan J. Tibshirani, and Larry Wasserman. Distribution-free predictive inference for regression. *Journal of the American Statistical Association*, 113(523): 1094–1111, 2018.
- Swee Kiat Lim, Aldrian Obaja Muis, Wei Lu, and Chen Hui Ong. MalwareTextDB: A database for annotated malware articles. In *Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1557–1567, Vancouver, Canada, 2017. Association for Computational Linguistics.
- Weitang Liu, Xiaoyun Wang, John D. Owens, and Yixuan Li. Energy-based out-of-distribution detection. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.
- Mike Mintz, Steven Bills, Rion Snow, and Daniel Jurafsky. Distant supervision for relation extraction without labeled data. In *Proceedings of the Joint Conference of the 47th Annual Meeting of the ACL and the 4th International Joint Conference on Natural Language Processing of the AFNLP*, pages 1003–1011, Suntec, Singapore, 2009. Association for Computational Linguistics.
- Alexander Ratner, Stephen H. Bach, Henry Ehrenberg, Jason Fries, Sen Wu, and Christopher Ré. Snorkel: Rapid training data creation with weak supervision. *Proceedings of the VLDB Endowment*, 11(3):269–282, 2017.
- Blake E. Strom, Andy Applebaum, Doug P. Miller, Kathryn C. Nickels, Adam G. Pennington, and Cody B. Thomas. MITRE ATT&CK: Design and philosophy. Technical Report MP 18-0944, MITRE Corporation, 2018.
- Vladimir Vovk, Alex Gammerman, and Glenn Shafer. *Algorithmic Learning in a Random World*. Springer, 2005.

## A Keyword List

The 21-term inclusion keyword list, matched case-insensitively against video title and description (Section 3), verbatim from `Scripts/fetch_transcripts.py`’s `KEYWORDS` constant:

```
ransomware, extortion, double extortion, encrypt, encryption,
decrypt, data leak, incident response, "ir", lockbit, alphv,
blackcat, conti, ryuk, cl0p, cl0p, akira, revil, sodinokibi,
ransomware-as-a-service, raas
```

The bare token “ir” is matched with surrounding spaces (" ir ") specifically to avoid matching it as a substring of unrelated words.

## NeurIPS Paper Checklist

### 1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [\[Yes\]](#)

Justification: The abstract and introduction (once finalized – see remaining \todo markers) state the corpus, the label-quality audit methodology and findings, weak-label baseline results explicitly framed as pipeline-agreement rather than verified accuracy, and OOD-robustness results as the paper’s contributions, matching Sections 3–8.

Guidelines:

- The answer [\[N/A\]](#) means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [\[No\]](#) or [\[N/A\]](#) answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

### 2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [\[Yes\]](#)

Justification: Section 10 opens with “No independent ground truth” as the paper’s central limitation rather than a minor bullet, and covers sampling bias, ASR transcription noise, the weak-label ceiling, and dual-use considerations.

Guidelines:

- The answer [\[N/A\]](#) means that the paper has no limitation while the answer [\[No\]](#) means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

### 3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper contains no theoretical results or proofs – it is an empirical dataset and baseline-framework contribution.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

#### 4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Section 3 specifies the channel/video selection criteria and full keyword list (Appendix A); Section 5 specifies the task taxonomy, feature representations, models, and CV protocols. All extraction, audit, and experiment code is in the public repository (see next item), including the exact hyperparameters used (default scikit-learn/LightGBM settings unless otherwise noted in code).

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
  - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
  - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
  - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).

- (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

## 5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: Dataset: <https://www.kaggle.com/datasets/henrykabuye/ransknow-v1> (CC0-1.0). Code: <https://github.com/Henrinnnes/RansKnow> (fetch pipeline, label-quality audit scripts, rk\_pipeline experiment framework, Phase 5 confound/counterfactual/perturbation scripts), released under the MIT License (LICENSE in the repository root).

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

## 6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: Section 5 specifies the three CV protocols (stratified random, channel-grouped, temporal) and the five-model/three-representation grid; exact hyperparameters (regularization, tree counts, TF-IDF vocabulary size, embedding model and chunking) are documented in Scripts/rk\_pipeline/ and referenced in Section 5.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

## 7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?



Answer: [No]

Justification: Reported macro-F1 values are single-seed, 5-fold (or single-split, for the temporal protocol) point estimates, not averaged or bounded across repeated random seeds.

**[TODO]** This is a real, addressable gap, not an inherent limitation – the framework’s split protocols already accept a seed parameter; re-running the Phase 1 grid across 3–5 seeds and reporting standard deviation is worth doing before submission, given how cheap most combinations were to run (see compute-resources item below).

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

## 8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Experiments ran on a local workstation: 32-core CPU (structured/TF-IDF representations, all sklearn/LightGBM models) plus 3×NVIDIA GTX 1070 Ti, 8GB each (embedding representation, sentence-transformer encoding). Observed wall-clock time: individual task/model/split combinations ranged from seconds to ~80s for the heaviest (RandomForest on 18-class multilabel TF-IDF); full grids (5 tasks × up to 5 models × 3 splits) completed in 20–45 minutes per representation.

**[TODO]** State total compute across the whole project, including exploratory runs that didn’t make the paper, per the guideline – not tracked precisely during development.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn’t make it into the paper).

## 9. Code of ethics



Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The authors have reviewed the NeurIPS Code of Ethics. The corpus is built from already-public YouTube content (Section 3); the current scope includes no human-subjects data collection (Section 4.3); dual-use considerations are discussed in Section 10.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

## 10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: Section 10's dual-use bullet discusses both the corpus's positive use (ransomware defense research, DFIR education) and the risk that structuring defender-oriented content could lower the effort to extract offense-relevant patterns.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

## 11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The dataset contains no operational malware, exploit code, or attack tooling – only transcripts of public educational/defensive security video content. No additional technical safeguards beyond the CC0 open license were implemented.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

## 12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: All source videos are publicly available on YouTube; per-video metadata (channel, title, URL, publish date) is retained in the dataset for attribution and traceability. The dataset itself is released under CC0-1.0 (Section 3; Kaggle listing [Kabuye, 2026]). The MITRE ATT&CK terminology this paper’s tactic taxonomy is aligned to is cited at first use (Strom et al., 2018, Section 1); ATT&CK is released by MITRE under a royalty-free license for research use with attribution (CC BY 4.0), which this citation satisfies.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, [paperswithcode.com/datasets](https://paperswithcode.com/datasets) has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset’s creators.

## 13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: RansKnow is a new dataset, documented via this paper (Sections 3–4) and the Kaggle dataset card. Machine-readable Croissant metadata, including the dataset-documentation fields (collection process, known biases, limitations, sensitive-information categories) required by the track’s hosting guidelines, is published at <https://github.com/Henrinnnes/RansKnow/blob/main/croissant.json> and validated with the `mlcroissant` validator (zero errors, zero warnings).

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.

- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

#### 14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: No crowdsourcing or human-subjects research is part of this paper’s current scope – the originally planned annotated evaluation set was descope (Section 4.3).

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

#### 15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: No human-subjects research was conducted (see previous item). Institutional ethics-approval constraints are discussed in Section 4.3 as the reason independent annotation was descope, which is distinct from having conducted such research.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

#### 16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [N/A]

Justification: LLM usage in this project does not constitute an important, original, or non-standard component of the core methods.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.

- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.